

COMP9313: Big Data Management



Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Chapter 12: Revision and Exam

Revision of Chapters Required in Exam

Topic 1 : MapReduce (Chapters 2-4)

Map and Reduce Functions

- Programmers specify two functions:
 - **map** (k_1, v_1) \rightarrow list [$\langle k_2, v_2 \rangle$]
 - ▶ Map transforms the input into key-value pairs to process
 - **reduce** ($k_2, [v_2]$) \rightarrow [$\langle k_3, v_3 \rangle$]
 - ▶ Reduce aggregates the list of values for each key
 - ▶ All values with the same key are sent to the same reducer
- Optionally, also:
 - **combine** ($k_2, [v_2]$) \rightarrow [$\langle k_3, v_3 \rangle$]
 - ▶ Mini-reducers that run in memory after the map phase
 - ▶ Used as an optimization to reduce network traffic
 - **partition** (k_2 , number of partitions) \rightarrow partition for k_2
 - ▶ Often a simple hash of the key, e.g., $\text{hash}(k_2) \bmod n$
 - ▶ Divides up key space for parallel reduce operations
 - **Grouping comparator**: controls which keys are grouped together for a single call to `Reducer.reduce()` function
- The execution framework handles everything else...

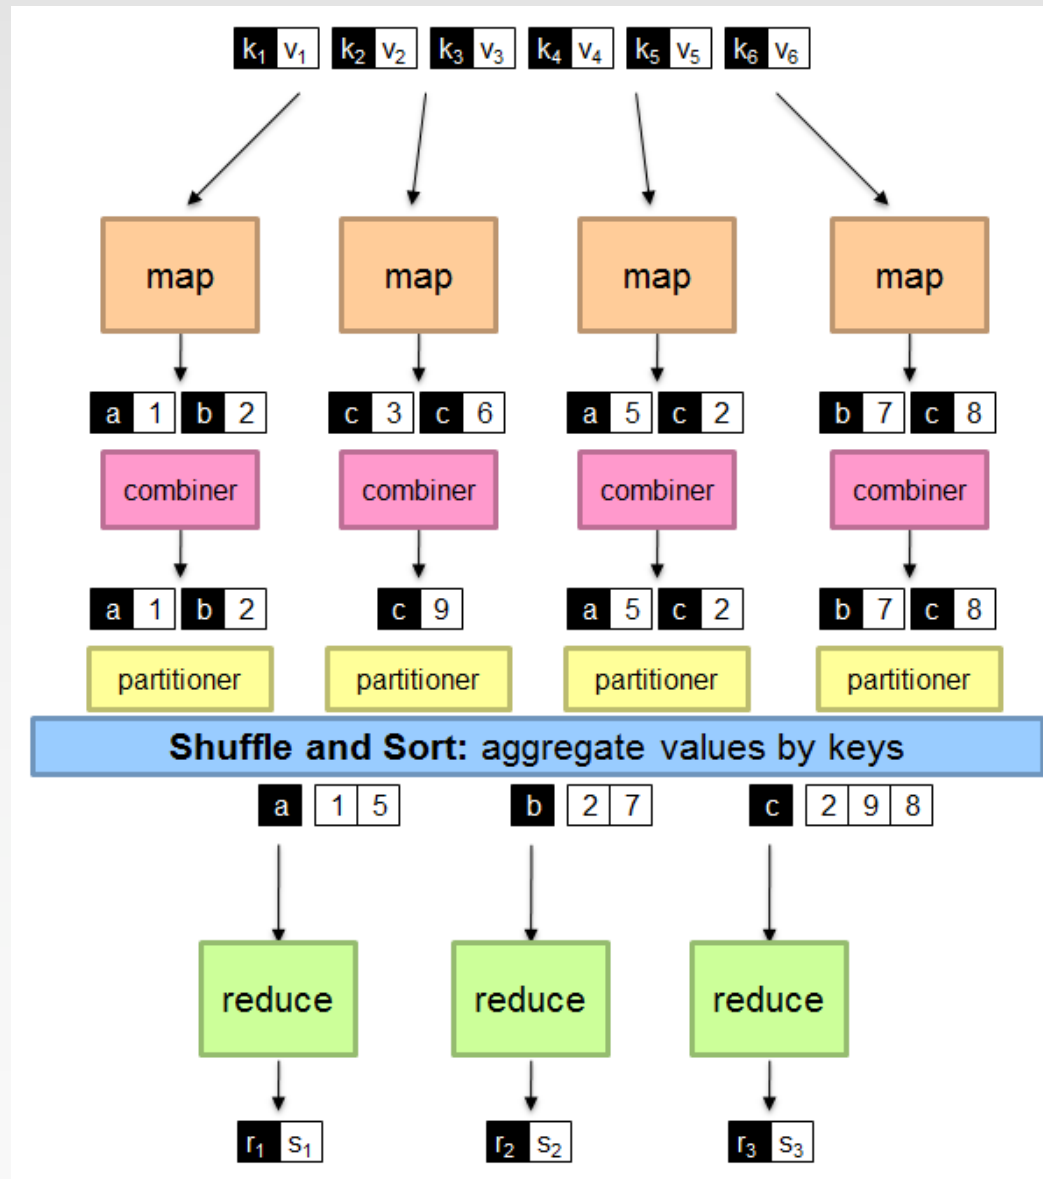
Combiners

- Often a Map task will produce many pairs of the form $(k, v_1), (k, v_2), \dots$ for the same key k
 - E.g., popular words in the word count example
- Combiners are a general mechanism to reduce the amount of intermediate data, thus saving network time
 - They could be thought of as “mini-reducers”
- Warning!
 - The use of combiners must be thought carefully
 - ▶ Optional in Hadoop: the correctness of the algorithm **cannot depend on** computation (or even execution) of the combiners
 - ▶ A combiner operates on each map output key. It must have the same output key-value types as the Mapper class.
 - ▶ A combiner can produce summary information from a large dataset because it replaces the original Map output
 - Works only if reduce function is commutative and associative
 - ▶ In general, reducer and combiner **are not interchangeable**

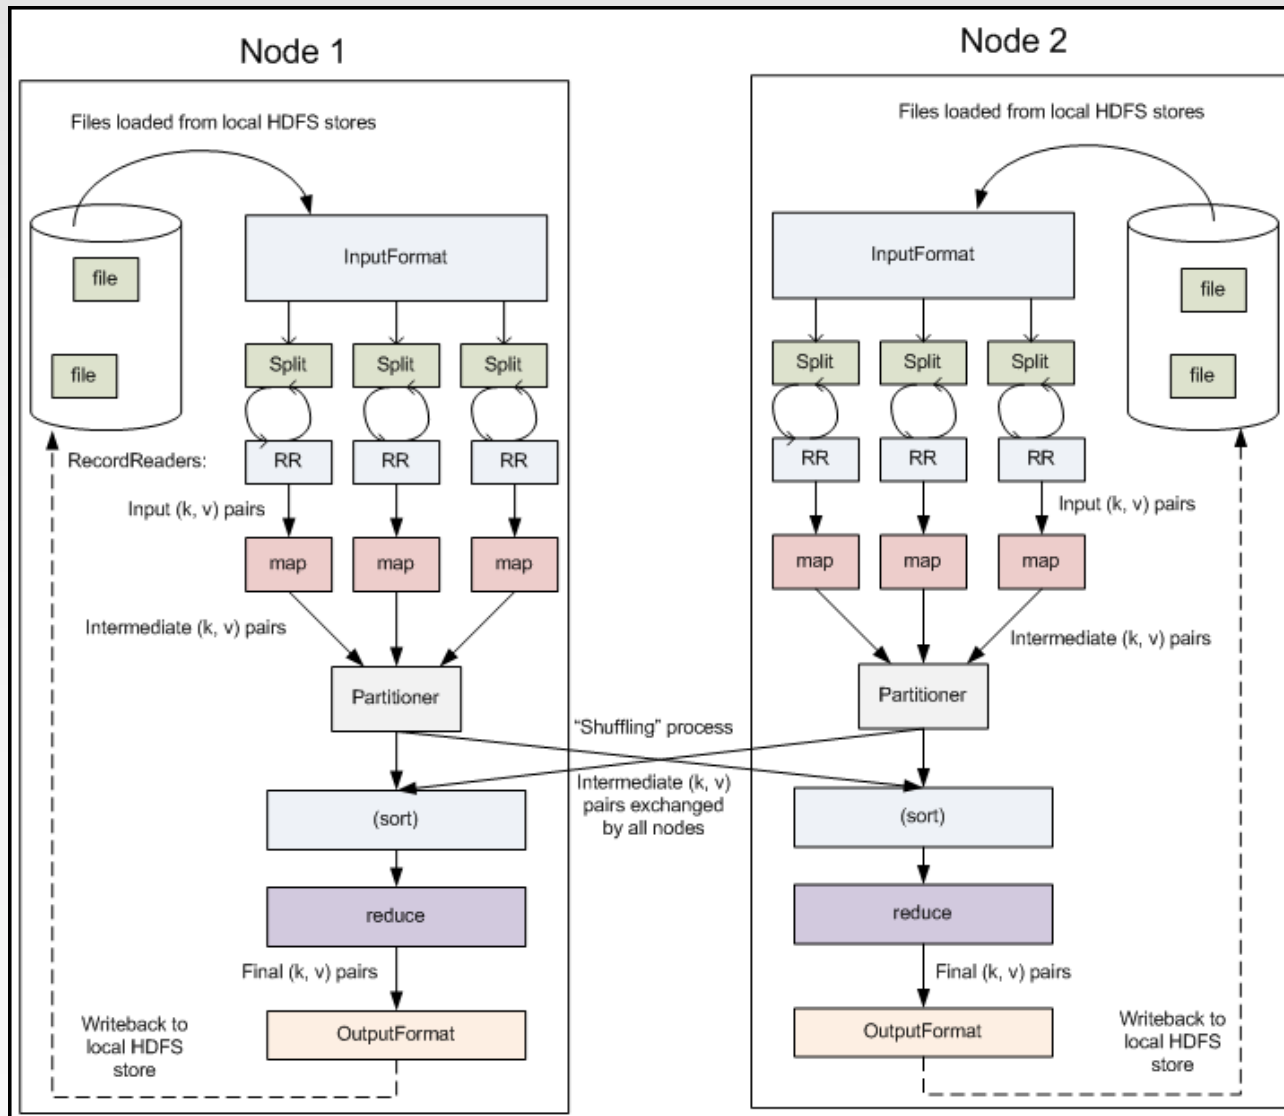
Partitioner

- Partitioner controls the partitioning of the keys of the intermediate map-outputs.
 - The key (or a subset of the key) is used to derive the partition, typically by a *hash function*.
 - The total number of **partitions** is the same as the number of reduce tasks for the job.
 - ▶ This controls which of the m reduce tasks the intermediate key (and hence the record) is sent to for reduction.
- System uses HashPartitioner by default:
 - $\text{hash}(\text{key}) \bmod R$
- Sometimes useful to override the hash function:
 - E.g., ***$\text{hash}(\text{hostname}(\text{URL})) \bmod R$*** ensures URLs from a host end up in the same output file
- Job sets Partitioner implementation (in Main)


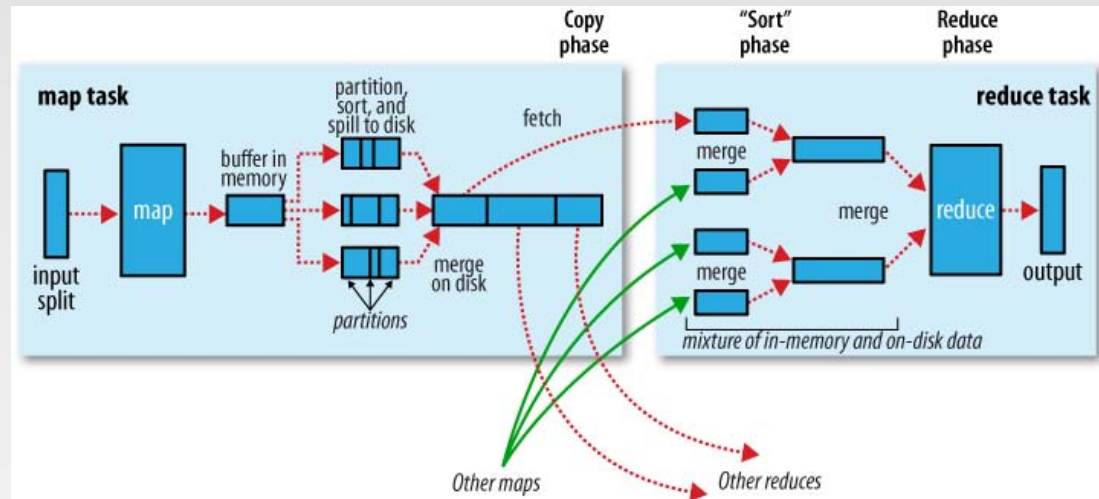
A Brief View of MapReduce



MapReduce Data Flow



MapReduce Data Flow

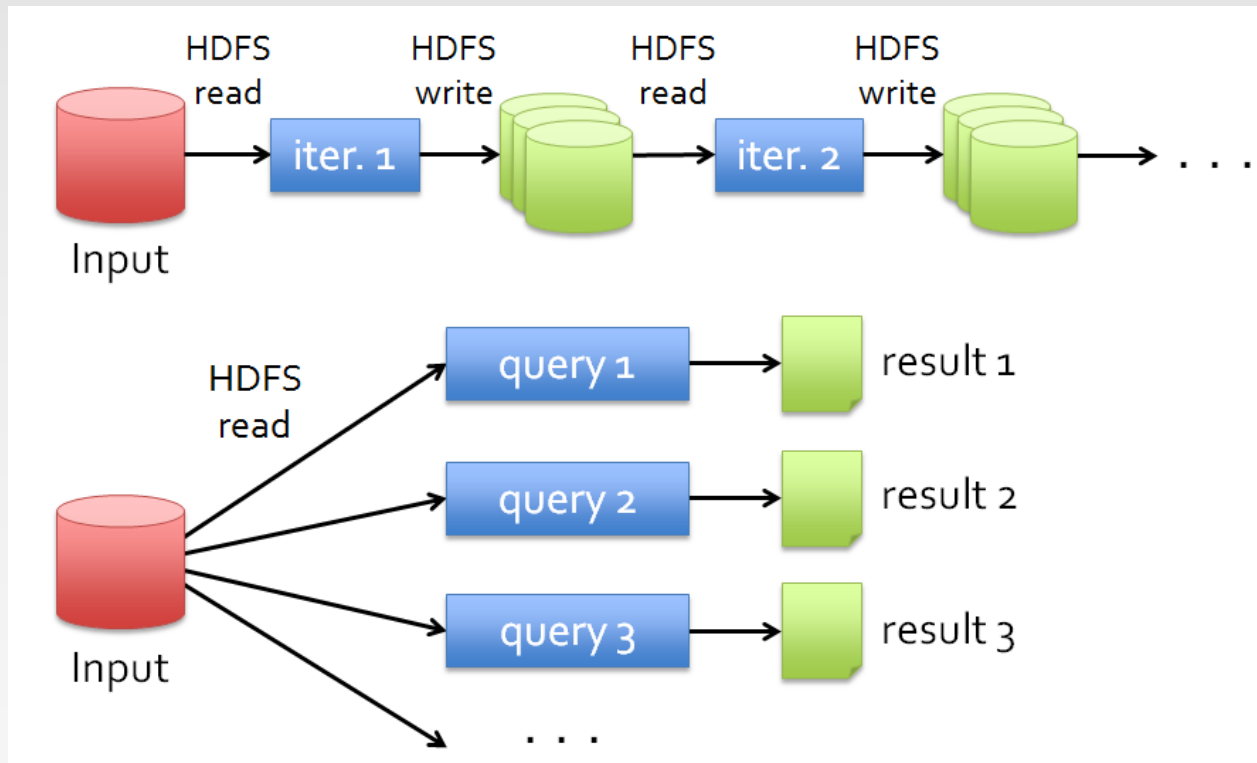


MapReduce Algorithm Design Patterns

- In-mapper combining, where the functionality of the combiner is moved into the mapper.
 - Scalability issue (**not suitable for huge data**) : More memory required for a mapper to store intermediate results
- The related patterns “pairs” and “stripes” for keeping track of joint events from a large number of observations.
- “Order inversion”, where the main idea is to convert the sequencing of computations into a sorting problem.
 - You need to guarantee that all key-value pairs relevant to the same term are sent to the same reducer
- “Value-to-key conversion”, which provides a scalable solution for secondary sorting.
 - Grouping comparator

Topic 2: Spark Core (Chapter 6)

Data Sharing in MapReduce




Slow due to replication, serialization, and disk IO

- Complex apps, streaming, and interactive queries all need one thing that MapReduce lacks:

Efficient primitives for **data sharing**

Data Sharing in Spark Using RDD

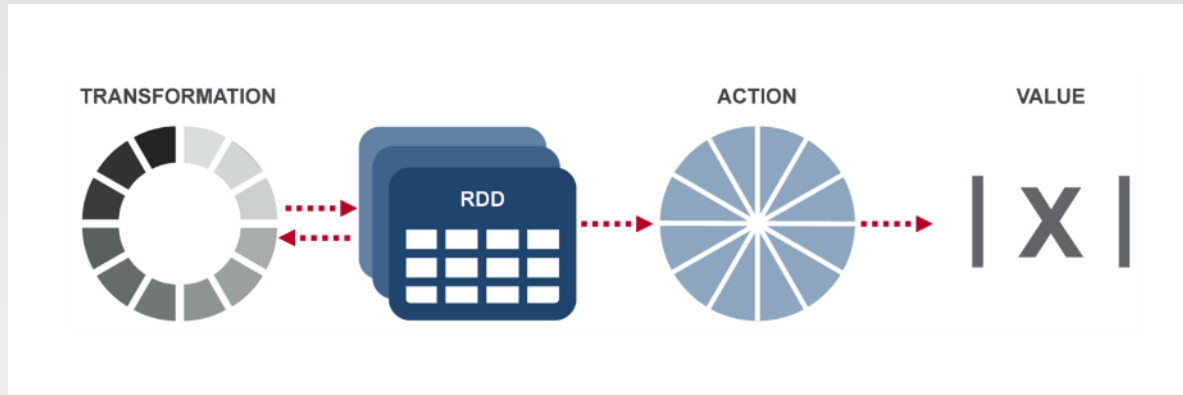


10-100 × faster than network and disk

What is RDD

- Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing. Matei Zaharia, et al. NSDI'12
 - RDD is a **distributed** memory abstraction that lets programmers perform **in-memory** computations on large clusters in a **fault-tolerant** manner.
- **Resilient**
 - Fault-tolerant, is able to recompute missing or damaged partitions due to node failures.
- **Distributed**
 - Data residing on multiple nodes in a cluster.
- **Dataset**
 - A collection of partitioned elements, e.g. tuples or other objects (that represent records of the data you work with).
- RDD is the primary data abstraction in Apache Spark and the core of Spark. It enables operations on collection of elements in parallel.

RDD Operations



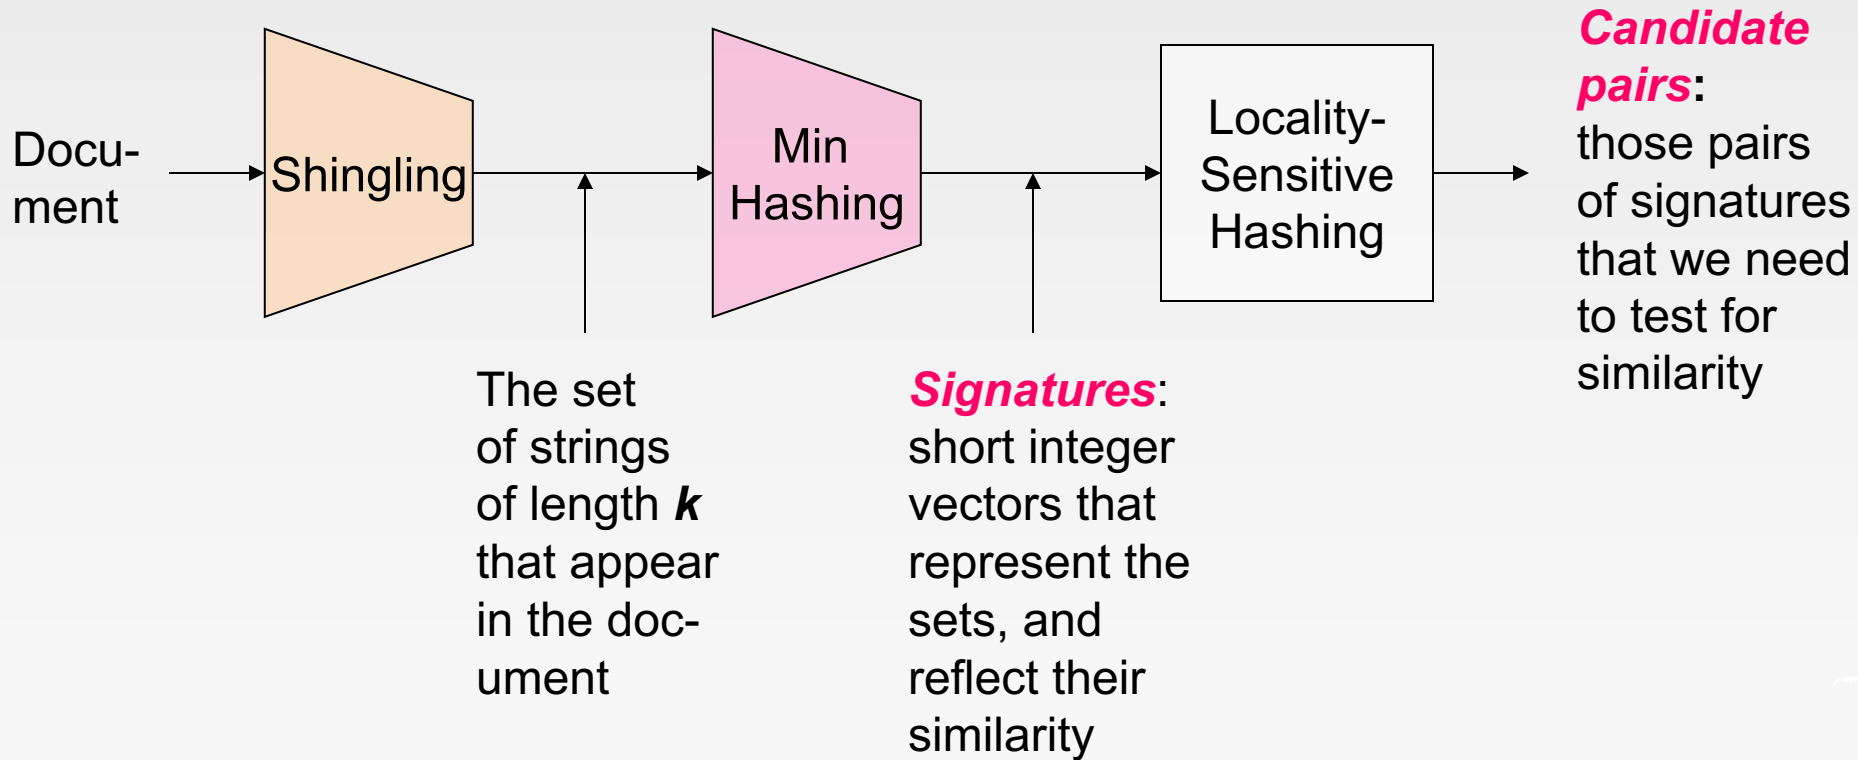
- **Transformation:** returns a new RDD.
 - Nothing gets evaluated when you call a Transformation function, it just takes an RDD and return a new RDD.
 - Transformation functions include *map*, *filter*, *flatMap*, *groupByKey*, *reduceByKey*, *aggregateByKey*, *filter*, *join*, *etc.*
- **Action:** evaluates and returns a new value.
 - When an Action function is called on a RDD object, all the data processing queries are computed at that time and the result value is returned.
 - Action operations include *reduce*, *collect*, *count*, *first*, *take*, *countByKey*, *foreach*, *saveAsTextFile*, *etc.*

RDD Operations

Transformations	$map(f : T \Rightarrow U) : RDD[T] \Rightarrow RDD[U]$ $filter(f : T \Rightarrow Bool) : RDD[T] \Rightarrow RDD[T]$ $flatMap(f : T \Rightarrow Seq[U]) : RDD[T] \Rightarrow RDD[U]$ $sample(fraction : Float) : RDD[T] \Rightarrow RDD[T]$ (Deterministic sampling) $groupByKey() : RDD[(K, V)] \Rightarrow RDD[(K, Seq[V])]$ $reduceByKey(f : (V, V) \Rightarrow V) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $union() : (RDD[T], RDD[T]) \Rightarrow RDD[T]$ $join() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (V, W))]$ $cogroup() : (RDD[(K, V)], RDD[(K, W)]) \Rightarrow RDD[(K, (Seq[V], Seq[W]))]$ $crossProduct() : (RDD[T], RDD[U]) \Rightarrow RDD[(T, U)]$ $mapValues(f : V \Rightarrow W) : RDD[(K, V)] \Rightarrow RDD[(K, W)]$ (Preserves partitioning) $sort(c : Comparator[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$ $partitionBy(p : Partitioner[K]) : RDD[(K, V)] \Rightarrow RDD[(K, V)]$
Actions	$count() : RDD[T] \Rightarrow Long$ $collect() : RDD[T] \Rightarrow Seq[T]$ $reduce(f : (T, T) \Rightarrow T) : RDD[T] \Rightarrow T$ $lookup(k : K) : RDD[(K, V)] \Rightarrow Seq[V]$ (On hash/range partitioned RDDs) $save(path : String) : \text{Outputs RDD to a storage system, e.g., HDFS}$

Topic 3: Finding Similar Items (Chapter 8)

■ The Big Picture



Shingling

- A **k -shingle** (or **k -gram**) for a document is a sequence of k tokens that appears in the doc
 - Tokens can be **characters**, **words** or something else, depending on the application
 - Assume tokens = characters for examples
- **Example:** $k=2$; document $D_1 = \text{abcab}$
Set of 2-shingles: $S(D_1) = \{\text{ab}, \text{bc}, \text{ca}\}$
- Documents that are intuitively similar will have many shingles in common.
 - **Example:** $k=3$, “The dog which chased the cat” versus “The dog that chased the cat”.
 - ▶ Only 3-shingles replaced are g_w, _wh, whi, hic, ich, ch_, and h_c.


Min-Hash Signatures

- Pick $K=100$ random permutations of the rows
- Think of $\text{sig}(\mathbf{C})$ as a column vector
- $\text{sig}(\mathbf{C})[i] =$ according to the i -th permutation, the index of the first row that has a 1 in column C


$$\text{sig}(\mathbf{C})[i] = \min (\pi_i(\mathbf{C}))$$

- **Note:** The sketch (signature) of document C is small ~ 100 bytes!
- **We achieved our goal!** We “compressed” long bit vectors into short signatures

Partition M into b Bands




Hashing Bands



b bands, r rows/band

- The probability that the minhash signatures for the documents agree in any one particular row of the signature matrix is t ($\text{sim}(C_1, C_2)$)
- Pick any band (r rows)
 - Prob. that all rows in band equal = t^r
 - Prob. that some row in band unequal = $1 - t^r$
- Prob. that no band identical = $(1 - t^r)^b$
- Prob. that at least 1 band identical = $1 - (1 - t^r)^b$

What b Bands of r Rows Gives You




Topic 4: Mining Data Streams (Chapter 9)

- Types of queries one wants on answer on a data stream: (we'll learn these today)
 - Sampling data from a stream
 - ▶ Construct a random sample
 - Queries over sliding windows
 - ▶ Number of items of type x in the last k elements of the stream
 - Filtering a data stream
 - ▶ Select elements with property x from the stream

Sampling Data Streams


- Since we can not store the entire stream, one obvious approach is to store a **sample**



- Two different problems:
 - (1) Sample a **fixed proportion** of elements in the stream (say 1 in 10)
 - ▶ As the stream grows the sample also gets bigger
 - (2) Maintain a **random sample of fixed size** over a potentially infinite stream
 - ▶ As the stream grows, the sample is of fixed size
 - ▶ At any “time” t we would like a random sample of s elements
 - What is the property of the sample we want to maintain?
For all time steps t , each of t elements seen so far has equal probability of being sampled

Fixup: DGIM Algorithm


- **Idea:** Instead of summarizing fixed-length blocks, summarize blocks with specific number of **1s**:
 - Let the block **sizes** (number of **1s**) increase exponentially
- When there are few 1s in the window, block sizes stay small, so errors are small



- Timestamps:
 - Each bit in the stream has a timestamp, starting from **1, 2, ...**
 - Record timestamps modulo **N (the window size)**, so we can represent any **relevant** timestamp in $O(\log_2 N)$ bits
 - ▶ E.g., given the windows size 40 (**N**), timestamp 123 will be recorded as 3, and thus the encoding is on 3 rather than 123

DGIM: Buckets

- A bucket in the DGIM method is a record consisting of:
 - (A) The timestamp of its end [$\mathcal{O}(\log N)$ bits]
 - (B) The number of 1s between its beginning and end [$\mathcal{O}(\log \log N)$ bits]
- Constraint on buckets:
 - Number of 1s must be a power of 2
 - That explains the $\mathcal{O}(\log \log N)$ in (B) above



Representing a Stream by Buckets

- The right end of a bucket is always a position with a 1
- Every position with a 1 is in some bucket
- Either **one** or **two** buckets with the same **power-of-2 number of 1s**
- Buckets do not overlap in timestamps
- **Buckets are sorted by size**
 - Earlier buckets are not smaller than later buckets
- Buckets disappear when their end-time is $> N$ time units in the past

Updating Buckets

- When a new bit comes in, drop the last (oldest) bucket if its end-time is prior to ***N*** time units before the current time
- 2 cases: Current bit is **0** or **1**
- If the current bit is 0: no other changes are needed
- If the current bit is 1:
 - (1) Create a new bucket of size 1, for just this bit
 - ▶ End timestamp = current time
 - (2) If there are now three buckets of size 1, combine the oldest two into a bucket of size 2
 - (3) If there are now three buckets of size 2, combine the oldest two into a bucket of size 4
 - (4) And so on ...

Example: Updating Buckets

Current state of the stream:

1001010110001011010101010101011010101010101110101010111010101110101000101100010

Bit of value 1 arrives

0010101100010110101010101010110101010101110101010111010101110101000101100101

Two white buckets get merged into a yellow bucket

00101011000101101010101010101101010101011101010101110101000101100101

Next bit 1 arrives, new orange white is created, then 0 comes, then 1:

010110001011010101010101011010101010111010101110101000101100101101

Buckets get merged...

010110001011010101010101011010101010111010101110101000101100101101

State of the buckets after merging

010110001011010101010101011010101010111010101110101000101100101101

Bloom Filter

- Consider: $|S| = m$, $|B| = n$
- Use k independent hash functions h_1, \dots, h_k
- Initialization:
 - Set B to all 0s
 - Hash each element $s \in S$ using each hash function h_i , set $B[h_i(s)] = 1$ (for each $i = 1, \dots, k$)
- Run-time:
 - When a stream element with key x arrives
 - ▶ If $B[h_i(x)] = 1$ for all $i = 1, \dots, k$ then declare that x is in S
 - That is, x hashes to a bucket set to 1 for every hash function $h_i(x)$
 - ▶ Otherwise discard the element x

Bloom Filter Example

- Consider a Bloom filter of size $m=10$ and number of hash functions $k=3$. Let $H(x)$ denote the result of the three hash functions.

- The 10-bit array is initialized as below

0	1	2	3	4	5	6	7	8	9
0	0	0	0	0	0	0	0	0	0

- Insert x_0 with $H(x_0) = \{1, 4, 9\}$

0	1	2	3	4	5	6	7	8	9
0	1	0	0	1	0	0	0	0	1

- Insert x_1 with $H(x_1) = \{4, 5, 8\}$

0	1	2	3	4	5	6	7	8	9
0	1	0	0	1	1	0	0	1	1

- Query y_0 with $H(y_0) = \{0, 4, 8\} \Rightarrow ???$

- Query y_1 with $H(y_1) = \{1, 5, 8\} \Rightarrow ???$ **False positive!**

- Another Example: <https://lilimlib.github.io/bloomfilter-tutorial/>

Topic 5: Recommender Systems (Chapter 11)

- Recommender systems
 - Content-based recommendation
 - Collaborative recommendation
 - ▶ User-user collaborative filtering
 - ▶ Item-item collaborative filtering
 - ~~Knowledge-based recommendation~~

	Avatar	LOTR	Matrix	Pirates
Alice	1		0.2	
Bob		0.5		0.3
Carol	0.2		1	
David				0.4

Final exam

- Final written exam (100 pts)
- Five questions in total on five topics
- Two hours
- Closed book exam
- If you are ill on the day of the exam, do not attend the exam – I will not accept any medical special consideration claims from people who already attempted the exam.

Exam Questions

- Question 1 MapReduce
 - Part A: MapReduce concepts
 - Part B: MapReduce algorithm design
- Question 2 Spark
 - Part A: Spark concepts
 - Part B: Show output of the given code
 - Part C: Spark algorithm design
- Question 3 Finding Similar Items
 - Shingling, Min Hashing, LSH
- Question 4 Mining Data Streams
 - Sampling, DGIM, Bloom filter
- Question 5 Recommender Systems

myExperience Survey

Give us a grade

UNSW has a new student course survey
– **myExperience**

Look out for your email invitation and
for links in Moodle

Fill out the survey to help us improve
your courses and teaching at UNSW

 **myExperience**

Thank you!